

This is a program that computes a hash value for an input file using multi-threading to improve the execution time. The program allows the user to input a file name and a desired number of threads to assist in the hashing of the file.


The experimental set-up was on a server which has the following specs: x86_64 architecture, with the Intel(R) Xeon(R) CPU E5-3695 v2 @ 2.40GHz CPU, made by GenuineIntel that consisted of 48 CPUs (2 sockets, 12 cores per socket, and 2 threads per core).

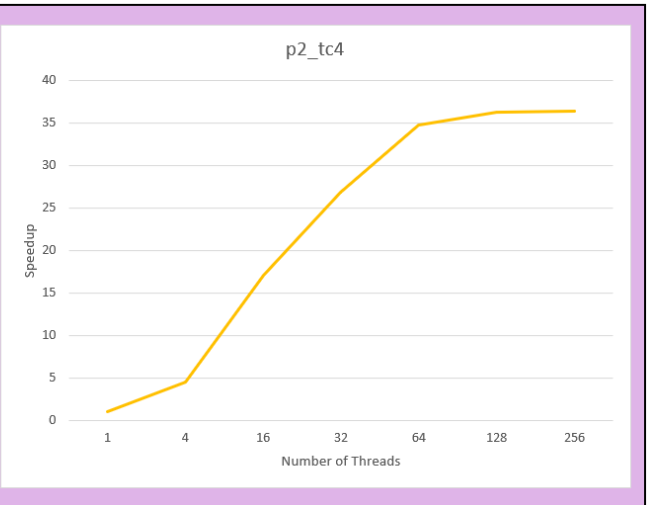
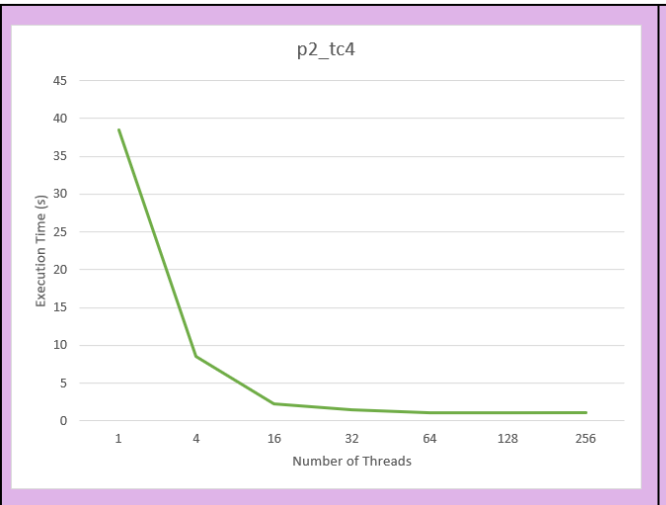
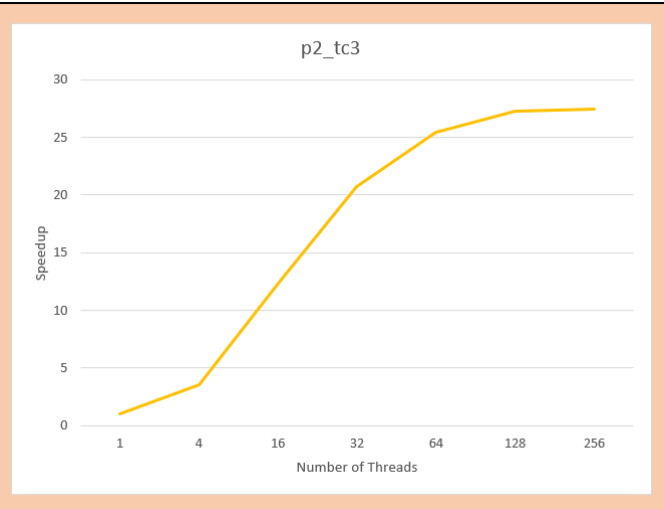


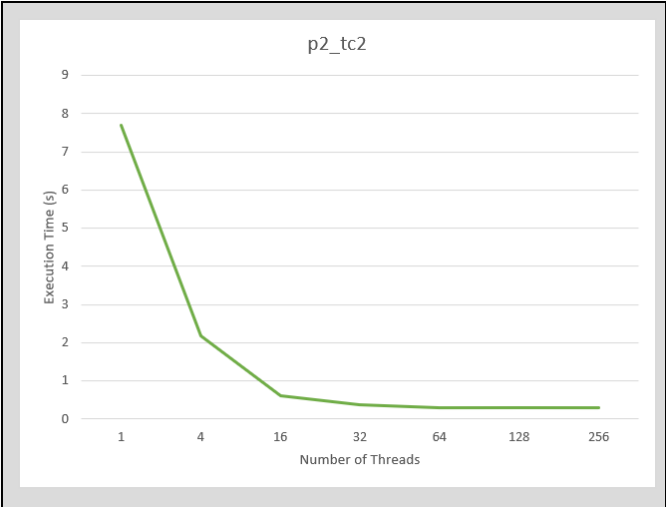
The main point of interest was to see how the runtime is affected as the number of threads is increased on the same file input. The execution of the program consisted of five test files: each file progressively bigger in size. The experiment tested it out at these increments of threads: 1, 4, 16, 32, 64, 128 and 256. Each increment we tested around 3-5 times and took the average of those runs to ensure an accurate execution time (I did not have independent use of the server where the test ran so other programs executing at the same time could cause ours to slow).

What was discovered was that the execution time does consistently go down as threads are increased, with the exception of the running time of threads between 128 and 256, in some test cases their execution times were extremely close and only marginally faster, but in some cases the time was marginally slower. Overall, the speedup shows diminishing returns as the number of threads increases. The biggest spike of improved execution time happens from 4 to 32 threads, with threads 1-16 seeing an approximate doubling of speedup (this is the only time in which the speedup is close to proportional to the number of threads increased), this slows down slightly from 32-128 threads and significantly less improved (in some cases worse) from 128-256 threads. Below is a table of each test case (labelled: p2_tc0-p2_tc4) with execution time in seconds and speedup (time taken for single thread/ time taken for n threads).

	Number of Threads	1	4	16	32	64	128	256
p2_tc0	Execution Time (s)	1.96	0.54	0.189	0.119	0.0987	0.093	0.09929
	Speedup	1	3.62963	10.37037	16.47059	19.85816	21.07527	19.74016
p2_tc1	Execution Time (s)	4.9	1.038	0.33	0.201	0.159	0.158	0.177
	Speedup	1	4.720617	14.84848	24.37811	30.81761	31.01266	27.68362
p2_tc2	Execution Time (s)	7.697	2.165	0.622	0.376	0.299	0.2906	0.296
	Speedup	1	3.555196	12.3746	20.47074	25.74247	26.48658	26.00338
p2_tc3	Execution Time (s)	14.9	4.218	1.204	0.719	0.585	0.546	0.543
	Speedup	1	3.53248	12.37542	20.72323	25.47009	27.28938	27.44015
p2_tc4	Execution Time (s)	38.5	8.479	2.254	1.435	1.108	1.062	1.059
	Speedup	1	4.54063	17.08075	26.82927	34.74729	36.25235	36.35505

Each of these test cases are graphed below to show the execution time vs. number of threads, as well as the speedup vs number of threads.





In conclusion, the speedup achieved by adding threads was a large improvement to the program's ability to hash files efficiently. Depending on the number of threads added (minimum 4, maximum 256) showed an improvement of a minimum of ~3.53x speedup and a maximum of ~36.35x speedup.